

*“Deep Learning Frameworks for Natural Language Processing: BERT and GPT-2”
Research Computing Center Workshop — Spring 2020
Jeffrey Tharsen tharsen@uchicago.edu*

This Guide and ipynb files :

https://github.com/rcc-uchicago/BERT-ELMo_tutorial_Spring2020

1. “This Grad Student Used a Neural Network to Write His Papers” (GPT-2)

<https://futurism.com/grad-student-neural-network-write-papers>

2. What are Transformers?

	GPT	Bert	Transformer-XL	XLNet	GPT-2	UniLM	XLNet	RoBERTa	Albert	CTRL	DistilBert	Bart	T5
Encoder stack	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
Decoder stack	✗	✗	≈	✗	≈	✗	≈	✗	✗	≈	✗	✓	✓

Not all models implement the Encoder-Decoder architecture; they are actually only becoming popular now. Transformer-XL, GPT2, XLNet and CTRL approximate a decoder stack during generation by using the hidden state of the previous state as the key & values of the attention module. Side note: all these 🍌 models are implemented in the [transformers library](#) or will be soon.

- <https://github.com/google-research/bert> (official Google repository for BERT)
- <https://medium.com/huggingface/encoder-decoders-in-transformers-a-hybrid-pre-trained-architecture-for-seq2seq-af4d7bf14bb8>
- <https://github.com/huggingface/transformers>

Model architectures

🍌 Transformers currently provides the following NLU/NLG architectures:

1. **BERT** (from Google) released with the paper [BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding](#) by Jacob Devlin, Ming-Wei Chang, Kenton Lee and Kristina Toutanova.
2. **GPT** (from OpenAI) released with the paper [Improving Language Understanding by Generative Pre-Training](#) by Alec Radford, Karthik Narasimhan, Tim Salimans and Ilya Sutskever.
3. **GPT-2** (from OpenAI) released with the paper [Language Models are Unsupervised Multitask Learners](#) by Alec Radford*, Jeffrey Wu*, Rewon Child, David Luan, Dario Amodei** and Ilya Sutskever**.
4. **Transformer-XL** (from Google/CMU) released with the paper [Transformer-XL: Attentive Language Models Beyond a Fixed-Length Context](#) by Zihang Dai*, Zhilin Yang*, Yiming Yang, Jaime Carbonell, Quoc V. Le, Ruslan Salakhutdinov.
5. **XLNet** (from Google/CMU) released with the paper [XLNet: Generalized Autoregressive Pretraining for Language Understanding](#) by Zhilin Yang*, Zihang Dai*, Yiming Yang, Jaime Carbonell, Ruslan Salakhutdinov, Quoc V. Le.
6. **XLNet** (from Facebook) released together with the paper [Cross-lingual Language Model Pretraining](#) by Guillaume Lample and Alexis Conneau.
7. **RoBERTa** (from Facebook), released together with the paper [A Robustly Optimized BERT Pretraining Approach](#) by Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, Veselin Stoyanov.
8. **DistilBERT** (from HuggingFace), released together with the paper [DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter](#) by Victor Sanh, Lysandre Debut and Thomas Wolf. The same method has been applied to compress GPT2 into **DistilGPT2**, RoBERTa into **DistilRoBERTa**, Multilingual BERT into **DistilmBERT** and a German version of DistilBERT.
9. **CTRL** (from Salesforce) released with the paper [CTRL: A Conditional Transformer Language Model for Controllable Generation](#) by Nishit Shrivastava*, Bryan McCann*, Lev R. Varshney, Caiming Xiong and Richard Socher.
10. **Camembert** (from Inria/Facebook/Sorbonne) released with the paper [Camembert: A Tasty French Language Model](#) by Louis Martin*, Benjamin Muller*, Pedro Javier Ortiz Suárez*, Yoann Dupont, Laurent Romary, Éric Villemonte de la Clergerie, Djamel Seddah and Benoît Sagot.
11. **ALBERT** (from Google Research and the Toyota Technological Institute at Chicago) released with the paper [ALBERT: A Lite BERT for Self-supervised Learning of Language Representations](#), by Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, Radu Soricut.
12. **T5** (from Google AI) released with the paper [Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer](#) by Colin Raffel and Noam Shazeer and Adam Roberts and Katherine Lee and Sharan Narang and Michael Matena and Yangqi Zhou and Wei Li and Peter J. Liu.
13. **XLNet-RoBERTa** (from Facebook AI), released together with the paper [Unsupervised Cross-lingual Representation Learning at Scale by Alexis Conneau*, Kartikay Khandelwal*, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Edouard Grave, Myle Ott, Luke Zettlemoyer and Veselin Stoyanov.](#)
14. **MBERT** (from Facebook), released together with the paper [A Super-sized Multilingual BERT for Classifying Images and Text](#) by Douwe Kiela, Suvrat Bhooshan, Hamed Firooz, Davide Testuggine.
15. **FlauBERT** (from CNRS) released with the paper [FlauBERT: Unsupervised Language Model Pre-training for French](#) by Hang Le, Loïc Vial, Jibril Frej, Vincent Segonne, Maximin Coavoux, Benjamin Lecouteux, Alexandre Allauzen, Benoît Crabbé, Laurent Besacier, Didier Schwab.
16. **BART** (from Facebook) released with the paper [BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension](#) by Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Ves Stoyanov and Luke Zettlemoyer.
17. **ELECTRA** (from Google Research/Stanford University) released with the paper [ELECTRA: Pre-training text encoders as discriminators rather than generators](#) by Kevin Clark, Minh-Thang Luong, Quoc V. Le, Christopher D. Manning.
18. **DialoGPT** (from Microsoft Research) released with the paper [DialoGPT: Large-Scale Generative Pre-training for Conversational Response Generation](#) by Yizhe Zhang, Siqi Sun, Michel Galley, Yen-Chun Chen, Chris Brockett, Xiang Gao, Jianfeng Gao, Jingjing Liu, Bill Dolan.
19. **Reformer** (from Google Research) released with the paper [Reformer: The Efficient Transformer](#) by Nikita Kitaev, Lukasz Kaiser, Anselm Levskaya.
20. **MarianMT** Machine translation models trained using OPUS data by Jörg Tiedemann. The Marian Framework is being developed by the Microsoft Translator Team.
21. **Longformer** (from AllenAI) released with the paper [Longformer: The Long-Document Transformer](#) by Iz Beltagy, Matthew E. Peters, Arman Cohan.
22. **Other community models**, contributed by the community.

These implementations have been tested on several datasets (see the example scripts) and should match the performances of the original implementations (e.g. ~93 F1 on SQUAD for BERT Whole-Word-Masking, ~88 F1 on RoCStories for OpenAI GPT, ~18.3 perplexity on WikiText-103 for Transformer-XL, ~0.916 Pearson R coefficient on STS-B for XLNet). You can find more details on the performances in the Examples section of the [documentation](#).

- <https://medium.com/@ngwaifoong92/beginners-guide-to-retrain-gpt-2-117m-to-generate-custom-text-content-8bb5363d8b7f>
- <http://jalamar.github.io/a-visual-guide-to-using-bert-for-the-first-time/> (DistilBert)

3. What Tasks can Transformers be trained for?

- <https://github.com/huggingface/transformers/blob/master/examples/README.md>

The Big Table of Tasks

Task	Example datasets	Trainer support	TFTrainer support	pytorch-lightning	Colab
language-modeling	Raw text	✓	-	-	 Open in Colab
text-classification	GLUE, XNLI	✓	✓	✓	 Open in Colab
token-classification	CoNLL NER	✓	✓	✓	-
multiple-choice	SWAG, RACE, ARC	✓	✓	-	 Open in Colab
question-answering	SQuAD	-	✓	-	-
text-generation	-	-	-	-	 Open in Colab
distillation	All	-	-	-	-
summarization	CNN/Daily Mail	-	-	-	-
translation	WMT	-	-	-	-
bertology	-	-	-	-	-
adversarial	HANS	-	-	-	-

- <https://towardsdatascience.com/bert-for-dummies-step-by-step-tutorial-fb90890ffe03>
- <https://medium.com/@Moscow25/the-best-deep-natural-language-papers-you-should-read-bert-gpt-2-and-looking-forward-1647f4438797>
- <https://medium.com/@ranko.mosaic/googles-bert-nlp-5b2bb1236d78>

4. Jupyter Notebook : Single-token prediction with Google's BERT (via Pytorch) **BERT_Predict_masked_token.ipynb**

5. Jupyter Notebook: Next sentence prediction with OpenAI's GPT-2 **BERT+GPT2_sentence_prediction.ipynb**

<https://openai.com/blog/better-language-models/>

The code comes from this tutorial :

https://colab.research.google.com/github/huggingface/blog/blob/master/notebooks/02_how_to_generate.ipynb

6. Training BERT for various tasks (on Google Cloud TPU via Colab)

Google Cloud TPU Quickstart:

<https://cloud.google.com/tpu/docs/quickstart>

<https://cloud.google.com/tpu/docs/setup-gcp-account>

In Cloud Console (“Dashboard”)

Create a cloud storage bucket; run the tutorial

Enable billing – you must put in a credit card number

“Start my free trial”

Create bucket (I used “bert_tutorial_12345”)

Colab : Predicting Movie Review Sentiment with BERT on TF Hub *(requires Tf 1.x, you may need to add !pip install --upgrade --force-reinstall tensorflow-gpu==1.15.0 at the beginning)*

[https://colab.research.google.com/github/google-](https://colab.research.google.com/github/google-research/bert/blob/master/predicting_movie_reviews_with_bert_on_tf_hub.ipynb)

[research/bert/blob/master/predicting_movie_reviews_with_bert_on_tf_hub.ipynb](https://colab.research.google.com/github/google-research/bert/blob/master/predicting_movie_reviews_with_bert_on_tf_hub.ipynb)

7. Training BERT for various tasks on Midway using pre-built models

A. Classification : MRPC (Sentence Equivalence) via GLUE

See Appendix I below

B. Question & Answer via SquAD (1.1 / 2.0)

<https://github.com/google-research/bert#squad-1.1>

<https://github.com/google-research/bert#squad-2.0>

8. How to build (“pre-train”) a BERT model from a custom source corpus on Midway

A. The basic procedure :

<https://github.com/google-research/bert#pre-training-with-bert>

B. Formatting the source corpus : *one line per sentence, one blank line between documents*

C. Creating vocab.txt

(WordPiece tokenization)

<https://github.com/kwonmha/bert-vocab-builder>

D. Setting up bert_config.json

(make sure the vocab length is correct)

Tips and best practices: <https://ufal.mff.cuni.cz/pbml/110/art-popel-bojar.pdf>

9. How to build a GPT-2 model from a custom corpus

<https://medium.com/@ngwaifoong92/beginners-guide-to-retrain-gpt-2-117m-to-generate-custom-text-content-8bb5363d8b7f>

<https://towardsdatascience.com/train-a-gpt-2-transformer-to-write-harry-potter-books-edf8b2e3f3db>

10. Retrieving and Visualizing the vectors from a BERT model (t-SNE with PCA)

A. Google’s method :

<https://github.com/google-research/bert#using-bert-to-extract-fixed-feature-vectors-like-elmo>

B. The “bert-embedding” library on Github :

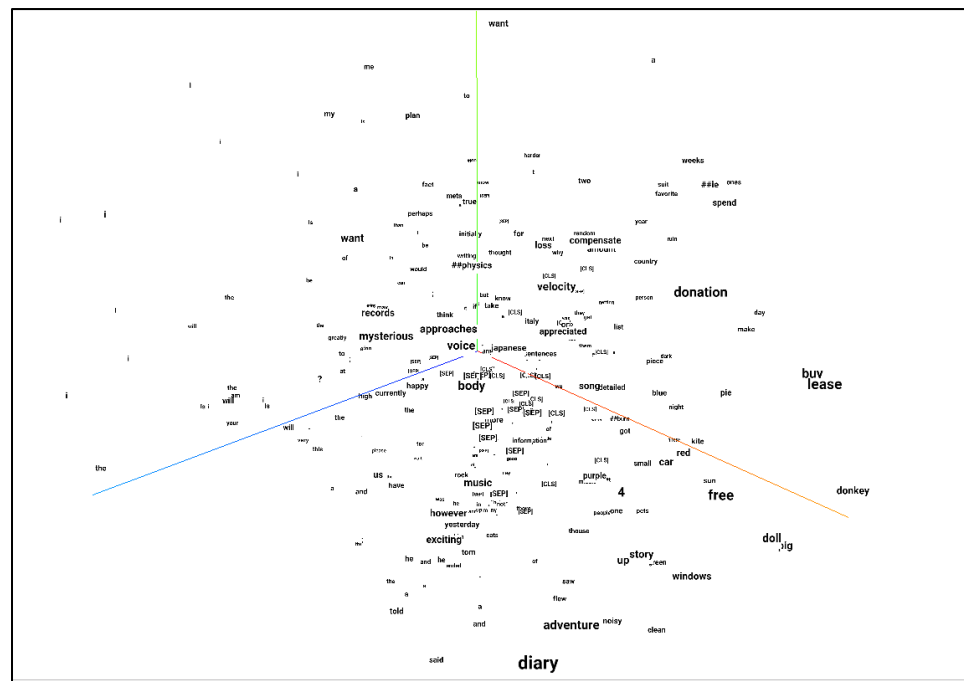
<https://github.com/negedng/bert-embedding>

C. Colab (for Tf 2.x) :

<https://colab.research.google.com/drive/1hMLd5-r82FrnFnBub-B-fVW78Px4KPX1#scrollTo=Iwew0KP8vRM>

D. Use Tensorflow's Tensorboard Embeddings Projector (3 dimensions) :

<https://towardsdatascience.com/bert-visualization-in-embedding-projector-dfe4c9e18ca9>



Appendix I. Using Midway to train BERT for Classification : MRPC (Sentence Equivalence) via GLUE

- I. Clone or download Google's main BERT repository:
git clone <https://github.com/google-research/bert.git>
- II. Download one of the pre-trained Google BERT models:
<https://github.com/google-research/bert#pre-trained-models>

e.g. **BERT-Base, Uncased** (recommended)

12-layer, 768-hidden, 12-heads, 110M parameters

https://storage.googleapis.com/bert_models/2018_10_18/uncased_L-12_H-768_A-12.zip

Open Firefox, download the zip file, then move it to your working directory and unzip it:

```
mv ~/Downloads/uncased_L-12_H-768_A-12.zip .
```

```
unzip uncased_L-12_H-768_A-12.zip
```

III. The model zip file will contain:

- | | |
|--|-------------------------------------|
| a. vocab.txt | BERT-Base, Uncased = 30,522 |
| b. bert_config.json | the all-important config file |
| c. bert_model.ckpt.meta | binary metadata file |
| d. bert_model.ckpt.index | binary index file |
| e. bert_model.ckpt.data-00000-of-00001 | the DATA, aka "the checkpoint file" |

IV. To run the GLUE test

First download download_glue_data.py script from

<https://gist.github.com/W4ngatang/>

module load Anaconda3/2018.12

Run : python download_glue_data.py

cd bert

Create an sinteractive session for yourself on Midway :

sinteractive -mem-per-cpu=24G

(or : sinteractive -mem-per-cpu=24G -partition=gpu2 -gres=gpu:1)

Load the modules and Tensorflow 1.13.1 environment :

module load Anaconda3/2018.12

source activate tf-cpu-1.13.1 (or source activate tf-gpu-1.13.1)

Export these directories (makes it easier to run run_classifier.py) :

export BERT_BASE_DIR=/path/to/bert/uncased_L-12_H-768_A-12

export GLUE_DIR=/path/to/glue_data

To train the classifier, run this:

```
python run_classifier.py \  
  --task_name=MRPC \  
  --do_train=true \  
  --do_eval=true \  
  --data_dir=$GLUE_DIR/MRPC \  
  --vocab_file=$BERT_BASE_DIR/vocab.txt \  
  --bert_config_file=$BERT_BASE_DIR/bert_config.json \  
  --init_checkpoint=$BERT_BASE_DIR/bert_model.ckpt \  
  --max_seq_length=128 \  
  --train_batch_size=32 \  
  --learning_rate=2e-5 \  
  --num_train_epochs=3.0 \  
  --output_dir=/tmp/mrpc_output/    → change this to somewhere useful
```

Then to run it, set TRAINED_CLASSIFIER to the output directory

```
export TRAINED_CLASSIFIER=/tmp/mrpc_output/    → what you used above
```

*Create the file test.tsv in the /bert directory (see below for a sample);
the process will create test_results.tsv in your output_dir.*

When test.tsv is ready, run this to create test_results.tsv in the output_dir :

```
python run_classifier.py \  
  --task_name=MRPC \  
  --do_predict=true \  
  --data_dir=$GLUE_DIR/MRPC \  
  --vocab_file=$BERT_BASE_DIR/vocab.txt \  
  --bert_config_file=$BERT_BASE_DIR/bert_config.json \  
  --init_checkpoint=$TRAINED_CLASSIFIER \  
  --max_seq_length=128 \  
  --output_dir=/tmp/mrpc_output/    → change this also to somewhere useful,  
                                     I use mrpc_test
```

Here are samples of both of these files (note that they are tab-delimited):

test.tsv :

guid	text_a	text_b
casd4	I hate pizza, as everone knows.	Pizza is the food I hate the most.
3ndf9	She is the greatest singer in the world.	As a singer, she's just the very best.
abcde	She loves me not.	Are penguins cold?

test_results.tsv will contain a list of vectors, to combine with test.tsv, do:

```
import pandas as pd #read the original test data for the text and id
df_test = pd.read_csv('bert/test.tsv', sep='\t')
#read the results data for the probabilities
df_result = pd.read_csv('mrpc_test/test_results.tsv', sep='\t', header=None) #change to the output dir
#create a new dataframe
df_map_result = pd.DataFrame({'guid': df_test['guid'],
                              'text_a': df_test['text_a'],
                              'text_b': df_test['text_b'],
                              'label': df_result.idxmax(axis=1)})
#view the first 3 rows of the newly created dataframe
print(df_map_result[0:3])
```

The final results should look something like this :

	guid	text_a	text_b	label
0	casd4	I hate pizza, as everone knows.	Pizza is the food I hate the most.	1
1	3ndf9	She is the greatest singer in the world.	As a singer, she's just the very best.	1
2	abcde	She loves me not.	Are penguins cold?	0